

# Chap 6 SOTA method

2022年2月2日 16:58

- 
- 两个 学派
  - Two lines of works on policy optimization:
    - ① Policy Gradient→Natural Policy Gradient/TRPO→ACKTR→PPO
    - ② Q-learning→DDPG→TD3→SAC
- 第一派：policy based
  - TRPO
    - 费舍尔信息矩阵
      - [Fisher Information Matrix - Agustinus Kristiadi's Blog](#): 主要证明了 费舍尔矩阵 可用于估计 黑塞矩阵。而黑塞矩阵常用于多元泰勒展开函数近似，这是计算TRPO参数更新的基础
      - 费舍尔信息矩阵 定义
        - ◆ 首先在极大似然估计原理中，假设有一个概率模型 $p(x|\theta)$  ( $\theta$ 是模型参数)，假设有大量采样 $x_i$ ，则要最大化 $p(x_i|\theta)$ ，一般是对这个函数求关于 $\theta$ 的极值，假设函数连续，极值处梯度一定为0
        - ◆ 因此定义评估函数 $s(\theta)$ 。 $\theta$ 是向量， $s$ 也是向量

$$s(\theta) = \nabla_{\theta} \log p(x|\theta),$$

- ◆ 区分两个概率：客观概率，和我们估计的主观概率模型 ( $p(x|\theta)$ )
- ◆ 本文章中，当使用主观概率采样时， $s(\theta)$ 的期望是0

$$\begin{aligned}\mathbb{E}_{p(x|\theta)} [s(\theta)] &= \mathbb{E}_{p(x|\theta)} [\nabla \log p(x|\theta)] \\&= \int \nabla \log p(x|\theta) p(x|\theta) dx \\&= \int \frac{\nabla p(x|\theta)}{p(x|\theta)} p(x|\theta) dx \\&= \int \nabla p(x|\theta) dx \\&= \nabla \int p(x|\theta) dx \\&= \nabla 1 \\&= 0\end{aligned}$$

而概率统计中学习的 极大似然估计法，是通过真实客观概率采样，当 $s(\theta)$ 的期望是0时， $\theta$ 就是真实参数的估计

- ◆ 费舍尔信息矩阵 为 $s(\theta)$  各分量 的协方差

$$\mathbb{E}_{p(x|\theta)} [(s(\theta) - 0)(s(\theta) - 0)^T].$$

$$\mathbf{F} = \mathbb{E}_{p(x|\theta)} [\nabla \log p(x|\theta) \nabla \log p(x|\theta)^T].$$

- 费舍尔矩阵等于负的黑塞矩阵的期望

- ◆ 黑塞矩阵分解

*Proof.* The Hessian of the log likelihood is given by the **Jacobian** of its gradient:

$$\begin{aligned} \mathbf{H}_{\log p(x|\theta)} &= \mathbf{J} \left( \frac{\nabla p(x|\theta)}{p(x|\theta)} \right) \\ &= \frac{\mathbf{H}_{p(x|\theta)} p(x|\theta) - \nabla p(x|\theta) \nabla p(x|\theta)^T}{p(x|\theta) p(x|\theta)} \\ &= \frac{\mathbf{H}_{p(x|\theta)} p(x|\theta)}{p(x|\theta) p(x|\theta)} - \frac{\nabla p(x|\theta) \nabla p(x|\theta)^T}{p(x|\theta) p(x|\theta)} \\ &= \frac{\mathbf{H}_{p(x|\theta)}}{p(x|\theta)} - \left( \frac{\nabla p(x|\theta)}{p(x|\theta)} \right) \left( \frac{\nabla p(x|\theta)}{p(x|\theta)} \right)^T, \end{aligned}$$

$$\begin{aligned} \mathbb{E}_{p(x|\theta)} [\mathbf{H}_{\log p(x|\theta)}] &= \mathbb{E}_{p(x|\theta)} \left[ \frac{\mathbf{H}_{p(x|\theta)}}{p(x|\theta)} - \left( \frac{\nabla p(x|\theta)}{p(x|\theta)} \right) \left( \frac{\nabla p(x|\theta)}{p(x|\theta)} \right)^T \right] \\ &= \mathbb{E}_{p(x|\theta)} \left[ \frac{\mathbf{H}_{p(x|\theta)}}{p(x|\theta)} \right] - \mathbb{E}_{p(x|\theta)} \left[ \left( \frac{\nabla p(x|\theta)}{p(x|\theta)} \right) \left( \frac{\nabla p(x|\theta)}{p(x|\theta)} \right)^T \right] \\ &= \int \frac{\mathbf{H}_{p(x|\theta)}}{p(x|\theta)} p(x|\theta) dx - \mathbb{E}_{p(x|\theta)} [\nabla \log p(x|\theta) \nabla \log p(x|\theta)^T] \\ &= \mathbf{H}_{\int p(x|\theta) dx} - \mathbf{F} \\ &= \mathbf{H}_1 - \mathbf{F} \\ &= -\mathbf{F}. \end{aligned}$$

- 自然梯度下降 Natural Gradient Descent

- [Natural Gradient Descent - Agustinus Kristiadi's Blog](#)
- Parameter Space versus Distribution Space
  - ◆ 一般梯度下降通过 欧氏范数 控制梯度下降步长。

$$\frac{-\nabla_{\theta} \mathcal{L}(\theta)}{\|\nabla_{\theta} \mathcal{L}(\theta)\|} = \lim_{\epsilon \rightarrow 0} \frac{1}{\epsilon} \arg \min_{d \text{ s.t. } \|d\| \leq \epsilon} \mathcal{L}(\theta + d).$$

- ◆ 自然梯度下降，根据更新前后概率模型分布的差别控制步长，具体是使用 KL 散度

- Importance Sampling

we can sample data from another distribution  $q$  and use the probability ratio between  $p$  and  $q$  to re-calibrate the result

□

$$\mathbb{E}_{x \sim p}[f(x)] = \int p(x) f(x) dx = \int q(x) \frac{p(x)}{q(x)} f(x) dx = \mathbb{E}_{x \sim q} \left[ \frac{p(x)}{q(x)} f(x) \right]$$

- 用在策略函数：

$$J(\theta) = \mathbb{E}_{a \sim \pi_\theta} [r(s, a)] = \mathbb{E}_{a \sim \hat{\pi}} \left[ \frac{\pi_\theta(s, a)}{\hat{\pi}(s, a)} r(s, a) \right]$$

## ▪ Trust Region Policy Optimization (TRPO)

- 结合了 IS 和 Natural Gradient

$$J_{\theta_{old}}(\theta) = \mathbb{E}_t \left[ \frac{\pi_\theta(a_t | s_t)}{\pi_{\theta_{old}}(a_t | s_t)} R_t \right]$$

$$\text{subject to } KL(\pi_{\theta_{old}}(\cdot | s_t) || \pi_\theta(\cdot | s_t)) \leq \delta$$

- 估计，使用多元函数的泰勒展开，并化简，

$$J_{\theta_t}(\theta) \approx g^T (\theta - \theta_t)$$

$$KL(\theta_t || \theta) \approx \frac{1}{2} (\theta - \theta_t)^T H (\theta - \theta_t)$$

where  $g = \nabla_\theta J_{\theta_t}(\theta)$  and  $H = \nabla_\theta^2 KL(\theta_t || \theta)$  and  $\theta_t$  is the old policy parameter

- 根据估计，写成

$$\theta_{t+1} = \arg \max_{\theta} g^T (\theta - \theta_t) \text{ s.t. } \frac{1}{2} (\theta - \theta_t)^T H (\theta - \theta_t) \leq \delta$$

- 解出参数更新的解析解

$$\theta_{t+1} = \theta_t + \sqrt{\frac{2\delta}{g^T H^{-1} g}} H^{-1} g$$

? 怎么算出来的

- 自然梯度

---

### Algorithm 1 Natural Policy Gradient

---

Input: initial policy parameters  $\theta_0$   
**for**  $k = 0, 1, 2, \dots$  **do**  
 Collect set of trajectories  $\mathcal{D}_k$  on policy  $\pi_k = \pi(\theta_k)$   
 Estimate advantages  $\hat{A}_t^{(k)}$  using any advantage estimation algorithm  
 Form sample estimates for  
 • policy gradient  $\hat{g}_k$  (using advantage estimates)  
 • and KL-divergence Hessian / Fisher Information Matrix  $\hat{H}_k$   
 Compute Natural Policy Gradient update:

$$\theta_{k+1} = \theta_k + \sqrt{\frac{2\delta}{\hat{g}_k^T \hat{H}_k^{-1} \hat{g}_k}} \hat{H}_k^{-1} \hat{g}_k$$

**end for**

---

① Sham Kakade. "A Natural Policy Gradient." NIPS 2001

- 计算H的逆，使用Conjugate Gradient迭代

- ① FIM and its inverse are very expensive to compute
- ② TRPO estimates the term  $x = H^{-1}g$  by solving the following linear equation  $Hx = g$
- ③ Consider the optimization for a quadratic equation

Solving  $Ax = b$  is equivalent to

$$x = \arg \max_x f(x) = \frac{1}{2}x^T Ax - b^T x$$

since  $f'(x) = Ax - b = 0$

- ④ Thus we can optimize the quadratic equation as

$$\min_x \frac{1}{2}x^T Hx - g^T x$$

- ⑤ Use **conjugate** gradient method to solve it. It is very similar to the gradient ascent but can be done in fewer iterations

□ TRPO总算法

- ① Resulting algorithm is a refined version of natural policy gradient

---

**Algorithm 3** Trust Region Policy Optimization

---

Input: initial policy parameters  $\theta_0$

**for**  $k = 0, 1, 2, \dots$  **do**

Collect set of trajectories  $\mathcal{D}_k$  on policy  $\pi_k = \pi(\theta_k)$

Estimate advantages  $\hat{A}_t^{\pi_k}$  using any advantage estimation algorithm

Form sample estimates for

- policy gradient  $\hat{g}_k$  (using advantage estimates)
- and KL-divergence Hessian-vector product function  $f(v) = \hat{H}_k v$

Use CG with  $n_{cg}$  iterations to obtain  $x_k \approx \hat{H}_k^{-1} \hat{g}_k$

Estimate proposed step  $\Delta_k \approx \sqrt{\frac{2\delta}{x_k^T \hat{H}_k x_k}} x_k$

Perform backtracking line search with exponential decay to obtain final update

$$\theta_{k+1} = \theta_k + \alpha^j \Delta_k$$

**end for**

---